
Recurrent neural network as a language model - technical report

Wojciech Zaremba
New York University

WOJ.ZAREMBA@GMAIL.COM

Abstract

Recurrent neural networks (RNN) offer a powerful framework to learn any arbitrary dependency. They are expressive as a finite memory Turing machine. However, their training is difficult and computationally expensive.

This technical note focuses on training RNNs for modeling language at the character level. We provide set of pragmatic recommendations about how to train a simple one layer RNN for this task. Moreover, we provide CPU and GPU Theano (Bergstra et al., 2010) code which reproduces results close to state-of-the-art on the Penn Treebank Corpus.

1. Introduction

Neural networks (NN) are stacked linear transformations alternated with non-linearities. Currently, state-of-the-art results in many computer vision tasks are achieved with feed-forward neural networks of this type. In feed-forward networks, computation flows in one direction from the input layer to the output layer(s). Recurrent neural networks (RNN) contain connections between instances of feed forward networks shifted in time. Such connections allow them to maintain *memory*, and perform prediction dependent on a history. Based on current advances in computer vision thanks to feed-forward networks, we are optimistic that models heavily utilizing RNNs can superior results to the current state-of-the-art on NLP tasks. Moreover, we believe that they might be crucial for further advances in computer vision (attention based models, and video prediction).

A common setting for RNNs is the prediction of the next element in a sequence. The input is a single element of a sequence, and a previous state. The network attempts at every stage to predict next element of sequence. We examine these models for language modelling and for simplicity, we

constrain ourselves to a character level language model.

The typical training procedure for RNNs is stochastic gradient descent (SGD). However, it is difficult to obtain effective RNN models by applying unconstrained SGD. Recurrency brings much higher expressive power compared to feed-forward networks, but also makes them more difficult to train. There are several well-known issues: (1) vanishing gradient; (2) exploding gradient and (3) short memory. We address exploding gradient issue by clipping gradients, but don't tackle the remaining problems.

We proceed by presenting related work (Section 2). Next, we describe our framework (Section 3), and finally we present experimental results (Section 4). Code reproducing experiments (and train any arbitrary RNN on CPU or GPU) is available online¹.

2. Related work

There has been extensive interest in different flavours of neural networks with recurrent connections (Hopfield, 1982; Hinton et al., 2006). These approaches consider recurrency not to account for time dependency, but for internal data dependency (e.g. correlation between pixel values).

In this note, we are mainly interested in RNNs which aim to predict a temporal sequence. (Mikolov, 2012) (Sutskever, 2013) consider training of such networks at the character and word level. Moreover, they analyze how best optimize such models (e.g. with Hessian-free method, or by clipping gradients).

(Graves, 2013) shows how the memory of the model may be extended by using Long-Short-Term-Memory units (LSTMs). Some evidence is shown that LSTMs prevent gradients from vanishing.

3. Framework

Our code is build in Theano. This Python framework permits the definition of symbolic expressions, and their automatic differentiation. Moreover, it compiles code to fast

¹<https://github.com/wojzaremba/rnn>

²<http://www.fit.vutbr.cz/~imikolov/rnnlm/simple-examples.tgz>

Generated sequence from input sequence "My name is"

| |
|---|
| My name is relatively ms. <unk> hill which i limit it social it can |
| My name is r. <unk> eurocom <unk> to <unk> more <unk> |
| My name is <unk> on smaller sales are clearance by long-term alterna |
| My name is reinvestment on senate is a voting oil co. his claim of s |
| My name issues culming was having a market acknowledged that it is in |

Table 3. Most of generated words are correct English words. Moreover, they are combined in an approximately grammatical way.

6. Acknowledgment

Thanks are due to Ilya Sutskever, Tomas Mikolov and Caglar Gulcehre for insightful discussions on how to train RNNs.

References

- Bergstra, James, Breuleux, Olivier, Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Desjardins, Guillaume, Turian, Joseph, Warde-Farley, David, and Bengio, Yoshua. Theano: a CPU and GPU math expression compiler. In *Proceedings of the Python for Scientific Computing Conference (SciPy)*, June 2010. Oral Presentation.
- Graves, Alex. Generating sequences with recurrent neural networks. *arXiv preprint arXiv:1308.0850*, 2013.
- Hinton, Geoffrey E, Osindero, Simon, and Teh, Yee-Whye. A fast learning algorithm for deep belief nets. *Neural computation*, 18(7):1527–1554, 2006.
- Hopfield, John J. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the national academy of sciences*, 79(8): 2554–2558, 1982.
- Mikolov, Tomáš. *Statistical language models based on neural networks*. PhD thesis, Ph. D. thesis, Brno University of Technology, 2012.
- Sutskever, Ilya. *Training Recurrent Neural Networks*. PhD thesis, University of Toronto, 2013.